

CloudVerse — Full Site Content

All 11 pages. Content only, no architecture.

PAGE 1: Homepage (cloudverse.ai)

Hero

H1: The Compute Economics Platform for the AI Era

Subhead: CloudVerse connects cloud infrastructure, data platforms, and AI workloads under one control plane. Cost context arrives where decisions get made, not after bills land.

CTA Primary: Book a Demo **CTA Secondary:** Explore the Platform

Stats strip

- 8+ GPU and LLM providers supported
 - Live cost-quality routing
 - Policy-bound at execution
 - \$738,983 recovered by a single customer
-

The problem (short, no header needed)

AI compute gets expensive by default. Data warehouse costs don't explain themselves. Cloud infrastructure spend grows faster than the teams managing it. The organizations that control this are the ones who ship cost intelligence into engineering workflows before the bill arrives, not after.

CloudVerse is how they do it.

Products section

H2: One control plane. Four compute surfaces.

AIX Routes every AI request to the lowest-cost model that meets your quality and latency requirements. Budget caps applied before spend occurs. Full decision trace on every request.

DataX Finds the queries running up your warehouse bill. Attributes cost to the query, pipeline, and team that ran it. Suggests the fix. Automates it when you're ready.

DevX Catches infrastructure cost regressions in pull requests before they reach production. Engineers see what their changes cost at the moment they can still change something.

FinOps Platform Multi-cloud cost intelligence across AWS, Azure, and GCP. Allocation, anomaly detection, commitment planning, and variance reporting on one model that reconciles to finance.

Case study: Berkshire Hathaway HomeServices

H2: How Berkshire Hathaway HomeServices recovered \$738,983

Stat 1: \$101,736 — annual recovery **Stat 2:** \$61,582 — single month recovery **Stat 3:** \$738,984 — total recovered

A growing mid-market AWS environment with fragmented tagging and no team-level attribution. CloudVerse connected spend to teams, surfaced the highest-impact anomalies, and gave the FinOps team a model that held up to finance review.

The waste was always there. It just had no address.

How it works (the lifecycle)

H2: Three stages. One control plane.

Inform Full cost attribution across cloud, data, and AI workloads. Down to the query, request, or deployment. Allocation that maps spend to the teams and systems that created it.

Optimize Recommendations with quantified expected savings. Cost-quality tradeoffs modelled before you commit. Engineering-grade signals, not finance reports.

Operate Automation paths with approval workflows, full audit logs, and rollback readiness. Budgets enforced at the routing layer. Governance that engineers don't route around.

Integration strip

H2: Connects to the stack your teams already use.

AWS, Azure, Google Cloud, Snowflake, Databricks, BigQuery, Microsoft Fabric, Azure Synapse, OpenAI, Anthropic, HuggingFace, Mistral AI, Groq, DeepSeek, GitHub, GitLab, Azure DevOps, Terraform, Pulumi, Kubernetes, and more.

Read-only by default. Automation is opt-in, scoped, and auditable.

Final CTA

H2: Connect your first account in under 30 minutes.

Most teams have their first non-obvious finding the same day.

CTA: Book a Demo **CTA Secondary:** Talk to Sales

PAGE 2: AIX (platform/aix)

Hero

H1: The AI Control Plane for Enterprise AI Compute

Subhead: Route, govern, audit, and optimize every AI request across models, clouds, GPUs, and private inference endpoints. Cost decisions happen at the routing layer, not in spreadsheets after the bill arrives.

CTA Primary: Book a Demo **CTA Secondary:** Explore the Platform

Stats strip

- 40-90% cost reduction across production workloads
 - Less than 15ms routing overhead added by AIX
 - 96.8% lower cost vs hardcoded Claude Sonnet (benchmarked)
 - 28.5% faster vs same hardcoded baseline
-

The problem

H2: AI compute is becoming too expensive to hardcode.

Most teams make AI infrastructure decisions once. Model endpoint, provider, region, failover logic, cost assumptions. These get buried in application code. Six months later, a cheaper model handles 70% of those requests at equivalent quality. Nobody knows.

Hardcoded decisions don't update when pricing shifts. They don't reroute when a provider degrades. They don't flag when a smaller model meets quality requirements at a third of the cost.

That gap is where most AI infrastructure spend goes unmanaged.

Before AIX (hardcoded)

- Model: Claude Sonnet
- Provider: Anthropic
- Region: us-east-1
- Routing: static
- Cost assumptions: set at launch, never revisited

After AIX (dynamic)

- Model: evaluated per request against constraints
 - Provider: scored across all connected providers
 - Region: applied per data residency policy
 - Routing: live signal evaluation on every request
 - Cost: attributed per request, per feature, per tenant
-

What AIX does

H2: How AIX controls every AI request.

AIX sits between your application and every AI provider you use. On every request, it evaluates all available routes against the constraints your team has defined, and returns the best one, with a fallback and a full decision log.

A gateway executes a routing rule you wrote. AIX evaluates whether that rule is still the right one.

Connect: Add provider API keys and endpoints. Minutes per provider. No application changes required.

Govern: Define constraints: latency ceiling, budget cap, allowed providers, compliance rules, quality floor. Per workload, per team, per region.

Score: Each route is evaluated on real-time latency distributions, effective unit cost, and availability signals. Live, not benchmarked.

Route: Policy constraints filter the scored routes. The best fit within policy is selected. Primary route and fallback returned together.

Measure: Cost per request, cost per feature, cost per tenant. Attribution automatic. No reconstruction from billing data.

Audit: Every routing decision logged: constraints active, routes evaluated, route selected, reason for selection, cost and latency outcome.

What AIX is not

H2: A gateway executes a rule you wrote. AIX decides what the rule should be.

Capability	Gateway / proxy	LLM gateway	Observability tool	AIX
Live multi-signal route scoring	No	Partial	No	Yes
Constraint-driven intent model	No	No	No	Yes
Cost-aware decisioning per request	No	Limited	Post-hoc	Yes
Policy enforcement at routing layer	Limited	Partial	No	Yes
Explainable routing decisions	No	No	No	Yes
Primary + fallback route output	Manual	Basic	No	Yes
GPU + managed API routing	No	No	No	Yes
Zero vendor lock-in	Partial	Partial	N/A	Yes

Observability tools tell you what happened after spend occurred. AIX acts before a request goes out. The routing decision, and the cost attached to it, is made explicitly before a token reaches any provider.

ROI and economics

H2: The cost of not routing.

Every hardcoded model endpoint is spending money without making a decision. AIX makes the decision explicitly. The saving is the difference.

At scale (monthly)

Monthly volume	Hardcoded spend	With AIX	Monthly saving
1M requests	\$2,980	\$298	\$2,682
5M requests	\$14,900	\$1,490	\$13,410
10M requests	\$29,800	\$2,980	\$26,820
50M requests	\$149,000	\$14,900	\$134,100

Assumptions: 40-90% cost reduction applied at 89% average. Hardcoded baseline is Claude Sonnet at \$0.00298/request.

Benchmark comparison (same workload) Without AIX: Claude Sonnet, 5,537ms latency, \$0.00298/req With AIX: GPT-4o-mini, 3,962ms latency, \$0.00010/req Result: 96.8% lower cost. 28.5% faster.

Three mechanisms

Cost arbitrage: AIX routes to the lowest-cost model that meets your quality floor. Prices change constantly. AIX adjusts automatically.

Latency wins: Faster routing reduces infrastructure overhead on time-sensitive workloads. Smaller models are often faster too.

Waste elimination: Stops premium model usage on tasks a smaller model handles equally well. Most applications have multiple task types. Few need the flagship model for all of them.

For FinOps teams

H2: From cloud cost reporting to AI compute economics.

AIX gives you cost-per-request, cost-per-feature, and cost-per-tenant allocation. Budget caps are enforced at the routing layer before spend occurs. No spreadsheet reconstruction after the fact.

- Cost-per-request, cost-per-feature, cost-per-tenant allocation — automatic
 - Budget caps enforced before a request goes out, not after the bill arrives
 - Scenario modelling: simulate cost impact before traffic or model changes go live
 - Anomaly detection: cost deviations flagged before they compound
 - Spend attribution: know which feature, tenant, or workload is driving AI cost
-

Enterprise governance

H2: Built for enterprise AI governance.

Every routing decision AIX makes is recorded, auditable, and explainable. Compliance controls are enforced at the routing layer, not bolted on after.

Multi-tenant isolation: Each team's workloads, policies, and cost data are separated.

Data residency controls: Route requests by region based on sovereignty requirements. EU, US, APAC per workload.

PII handling rules: PII detection enforced before provider selection. Sensitive requests never reach unapproved endpoints.

Budget caps: Hard spend ceilings applied before a request goes out.

Org/team policy scopes: Different teams run under different constraint sets. One platform, multiple policies.

Full execution trace logs: Every decision logged with constraints, candidates, selection, outcome.

No black boxes

H2: No black boxes.

Every routing decision includes an evidence summary: constraints active, routes evaluated, route selected, reason for selection, cost and latency outcome. Incident response, finance reviews, and compliance audits all work from the same traceable record.

```
json
{
  "requestId": "req_a1b2c3",
  "constraints": {
    "budget_cap": "$0.0005",
    "region": "EU",
    "latency_max": 4000
  },
  "routesEvaluated": 4,
  "selectedRoute": "GPT-4o-mini (Together AI, eu-west)",
  "reason": "Lowest cost within latency and region constraints",
  "result": {
    "cost": "$0.00010/req",
    "latency": "3,962ms",
    "fallback": "Mistral-Large (eu-central)"
  }
}
```

SOC2 and ISO-aligned architecture. Audit trails, access controls, and policy enforcement structured to support compliance documentation. Specific certification status confirmed during your evaluation.

Use cases

H2: Four problems AIX fixes.

1. Multi-provider cost arbitrage

Situation: A SaaS company runs 15M AI requests per month through a single hardcoded premium model.

The problem: \$44,700/month. 70% of those requests could be handled by a smaller model at equivalent quality.

How AIX solves it: Profiles each request type by task complexity. Routes to the lowest-cost model that meets the defined quality floor. No prompt changes. No application rewrites.

After AIX: ~\$12,600/month. 30% premium, 70% cost-optimised. Full cost allocation by request type.

2. Compliance-bound routing

Situation: A financial services firm operates in the EU, UK, and Singapore, each with different data residency requirements. Provider allowlists are maintained in a spreadsheet.

The problem: Manual enforcement. Unapproved providers get used in testing. Compliance audits require manual log reconstruction.

How AIX solves it: Residency rules configured per org and region in the AIX policy engine. PII handling rules enforced before provider selection. Every routing decision produces a machine-generated audit trace.

After AIX: Zero unapproved provider incidents post-deployment. Audit logs complete, not manually assembled. Security team removes spreadsheet enforcement process entirely.

3. GPU cost control for AI agents

Situation: An enterprise team runs AI agents for internal automation. Monthly AI spend varies by up to 3x depending on agent activity.

The problem: No cost visibility per agent run. Budget caps exist in documentation but are not enforced at the infrastructure layer. Spend spikes get raised by finance two weeks after the fact.

How AIX solves it: Per-agent budget constraints applied at the routing layer. When a ceiling is approached, AIX routes to lower-cost alternatives automatically. Cost per agent run visible in real time.

After AIX: Monthly spend variance drops from 3x to under 15%. Finance gets allocation data without engineering involvement. Budget enforcement becomes automated, not a policy conversation.

4. Multi-model AI product infrastructure

Situation: A product team builds an AI-native app with five distinct features, each with different latency, quality, and cost requirements. One model is hardcoded per feature.

The problem: Six months in, two routing choices are suboptimal. One due to pricing, one due to latency degradation. Fixing either requires a deployment cycle.

How AIX solves it: Each feature is a separate AIX workload with its own constraint profile. When a better-fit model becomes available, AIX routes there automatically. Routing updates happen at the constraint layer, not the code layer.

After AIX: Product team responds to model market changes in hours, not sprint cycles. Per-feature cost visible in the AIX dashboard. No deployment needed when providers or prices change.

Integrations

H2: Connect once. Route everywhere.

Add your providers once. AIX handles routing, fallback, and cost tracking across all of them. No code changes when you add a new provider.

Supported model providers

Provider	Models and notes
OpenAI	GPT-4o, GPT-4o-mini, GPT-4 Turbo, GPT-3.5 Turbo. Managed API.
Anthropic	Claude 3.5 Sonnet, Claude 3 Haiku, Claude 3 Opus. Managed API.
HuggingFace	Hosted inference endpoints and open models via HuggingFace Hub.
Mistral AI	Mistral Large, Mistral Small, Mixtral 8x7B. Managed API and self-hosted.
Llama / Ollama	Meta Llama 3 family via Ollama and compatible deployments.
Cohere	Command R, Command R+, reranking endpoints. Managed API.
Together AI	Open model inference at scale. Llama, Mistral, Qwen variants.
Groq	Ultra-low latency LPU inference. Llama and Mistral family.
DeepSeek	DeepSeek-V3, DeepSeek-R1. Managed API.

Private deployments and GPU infrastructure

AIX treats private GPU capacity as a first-class routing target alongside managed APIs. If you run models on dedicated hardware or a NeoCloud provider, AIX routes to them with the same cost and policy logic.

- Private model deployments: vLLM, TGI, custom inference
- GPU providers: CoreWeave, Lambda Labs, RunPod
- On-premises infrastructure: your own inference hardware

Why AIX

H2: Most tools route traffic. AIX decides where it should go.

A gateway executes whatever rule you gave it. If you told it to send traffic to Provider A, that is where it goes, even when Provider B is 90% cheaper and just as capable. A gateway has no scoring engine and no live signal evaluation. It is plumbing.

AIX is not a gateway replacement. It is the decision layer that makes routing intelligent rather than static.

AIX is the right fit when:

- You run AI workloads in production across more than one provider
- AI spend is material enough that optimisation has measurable impact
- Compliance requirements constrain provider or region selection
- Your team hardcoded model endpoints and has not revisited them

- You want cost-aware routing without rebuilding your application
-

Final CTA

H2: See AIX route your AI workloads end to end.

Connect your first account in under 30 minutes. Most teams have their first non-obvious finding the same day.

CTA: Get a Demo **CTA Secondary:** Talk to Sales

PAGE 3: DataX (platform/datax)

Hero

H1: Find the Queries Quietly Running Up Your Bill

Subhead: Query, dashboard, and dbt-model-level attribution across Snowflake, Databricks, BigQuery, Microsoft Fabric, and Synapse. Safe automation when you want it.

CTA Primary: Book a Demo **CTA Secondary:** Explore the Platform

Stats strip

- \$117.16 — full scan detected on a single query pattern
 - 334.6 GB — scanned per query
 - 77 runs at \$1.52 average cost
 - 6 warehouses supported
 - Read-only by default
-

The problem

H2: Your data warehouse is a blank check.

Most data teams find out about expensive queries the same way. The monthly bill arrives, it is higher than last month, and everyone starts guessing. The data is somewhere in the query logs. Finding it takes hours. Fixing it takes longer.

DataX attributes every dollar of warehouse spend to the query, pipeline, dashboard, and team that ran it. Not to "the data team." To the specific SELECT that ran 334.6 GB without a partition filter.

Stop waiting for the monthly bill to see who burned the budget.

The \$117 finding

H2: We find the leaks billing dashboards miss.

This is a real DataX finding. Not a mock. Not an illustration.

The \$117 full scan: A BigQuery SELECT scanning 334.6 GB per query due to missing partition pruning. Zero cache hit rate. \$117.16 in real cost. Fix: one-click partition pruning.

The pattern view billing never shows: DataX groups queries into cost-amplifying patterns, not one-off executions. This query ran 77 times at \$1.52 average cost for \$117.16 total. 100% scan rate. Tagged full-scan, spiky, fan-out. This is where hidden spend actually lives.

The high-frequency amplifier: A query running 77 times a month does not look expensive until the repetition multiplies the cost. DataX detects this automatically. Exact SQL available. Frequency plus cost math explained. Cache and materialisation fixes suggested.

What data teams unlock

H2: Warehouse cost intelligence, not just dashboards.

Query attribution Every query tied to the user, role, dashboard, model, or job that ran it. When the data team gets blamed, they can show exactly which workload and which team owns the cost.

Pattern detection Repeated expensive patterns surfaced with rewrite suggestions. The problem is usually a handful of query patterns running hundreds of times, not one catastrophic scan. DataX finds the pattern, not just the instance.

Predictive signals Detect warehouse cost spikes from queue depth and pattern shifts before they hit the bill. Most regressions are visible in telemetry before they become a finance conversation. DataX flags them before they do.

Safe automation One-click partition, cluster, and right-size fixes. Fully policy-bound and auditable. Automation in DataX controls when recommendations may be applied, and provides a complete audit trail of every decision and action. Opt-in, scoped, reversible.

Pricing philosophy

H2: DataX does not profit from your inefficiency.

DataX prices on the structural drivers of your data platform cost, not on billing noise.

What pricing reflects:

- Analytics compute footprint: warehouses, clusters, slots, capacities

- Analytics storage footprint: tables, partitions, datasets

What does not affect pricing:

- Sudden cost spikes or anomalies
- Inefficient queries or pipelines
- Retries, backfills, or bad deployments
- Alert volume or rule execution
- Short-term workload regressions

These are precisely the problems DataX is designed to detect and prevent. We do not monetize them.

As your platform becomes more efficient, your cost per unit of work improves without being penalised for growth.

Safe automation

H2: Automation without losing control.

DataX applies approved optimisations within policies you define, and logs every action for audit and rollback.

The automation model is about DataX behaviour, not your pipelines.

- Opt-in policies and scoped permissions
- Confidence scoring gates execution
- Approval workflows for higher-risk changes
- Every action logged with reason, impact, rollback readiness

Automation modes:

- Observe only
- Recommend only
- Auto-apply safe
- Approval required

Example audit event: Detected: Warehouse Oversizing (Snowflake). Reason: 95th percentile utilisation under 15% for 7 days. Action taken: Resized Medium to Small. Approved by: Auto-Apply (Policy: Safe Optimisation).

If it was not logged, it did not happen.

Supported platforms

H2: Enterprise integrations built for least privilege.

Connect platforms using read-only access by default. Enable automation only when you are ready: scoped, auditable, reversible.

Platform	What DataX covers
Snowflake	Credits, spillage, clustering, caching
Databricks	DBUs, clusters, shuffle behaviour
BigQuery	Scan economics, slots, reservations
Microsoft Fabric	CU debt, smoothing, throttling
Azure Synapse	DWUs, serverless TB scanned

Connection model:

1. Intent-based onboarding per platform type
2. Read-only permissions by default
3. Optional automation permissions — scoped, explicit
4. Metadata Pulse verifies telemetry before ingest
5. Least-privilege scripts your security team can audit

Read-only means read-only. DataX ingests metadata, query logs, and billing telemetry. It never touches your underlying data, workload code, or runtime configuration unless you explicitly grant automation permissions.

FAQ

Does DataX read our actual data? No. DataX reads query history and metadata only, never row contents. Read-only by default.

Which warehouses are supported? Snowflake, Databricks, BigQuery, Microsoft Fabric, Azure Synapse.

How does it work with dbt? DataX attributes costs to dbt model runs and flags high-cost models with optimisation suggestions. The attribution follows the DAG.

Can data engineers act on recommendations directly? Yes. One-click optimisation. All actions are policy-bound and auditable. No finance approval cycle required for routine fixes.

How does DataX pricing work? Pricing is calculated automatically after login based on your connected platforms. It reflects your steady-state analytics footprint, not spike volume or alert

activity.

Final CTA

H2: See DataX find a \$117 query in your warehouse.

Connect your first account in under 30 minutes. Most teams have their first non-obvious finding the same day.

CTA: Book a Demo **CTA Secondary:** Talk to Sales

PAGE 4: DevX (platform/devx)

Hero

H1: Catch Cost Regressions Before They Reach Production

Subhead: Every infrastructure PR gets an inline cost impact comment before reviewers see it. Engineers see what their changes cost at the moment they can still change something.

CTA Primary: Book a Demo **CTA Secondary:** See a Real PR Example

Stats strip

- +\$2.4k to \$2.9k flagged on a single PR
 - Free tier available
 - GitHub, GitLab, Azure DevOps native integrations
 - 7+ IaC formats supported
-

The problem

H2: Infrastructure cost mistakes do not look like mistakes.

Infrastructure decisions that cause cost regressions look like correct Terraform. The NAT gateway goes in because someone needed it for one sprint. Nobody removes it. The always-on compute instance gets sized for peak load. Load normalises. The instance stays.

Cost reports land three weeks after the deployment. The engineer who wrote the change has moved on to four other things. Nobody changes anything.

DevX puts a cost estimate on every PR before it merges. Engineers see impact at the moment they still have context on what they built and why.

Live PR example

H2: This is what a DevX PR comment looks like.

This is a real DevX output. An infrastructure change that looked routine. NAT gateway flag and a compute resize.

```
diff
resource "aws_nat_gateway" "this" {
-   count = 1
+   count = var.enable_nat_gateway ? 1 : 0
    allocation_id = aws_eip.nat.id
    subnet_id     = aws_subnet.public.id
}

resource "aws_instance" "app" {
-   instance_type = "t3.large"
+   instance_type = "t3.medium"
}
```

DevX cost estimate: +\$1.1k to \$1.4k per month

Primary drivers:

- NAT Gateway hourly + data processing charges
- Over-provisioned compute for observed utilisation

Why this matters:

- Always-on NAT in non-prod is a recurring cost with no production benefit
- Instance size exceeds observed utilisation

Suggested fix:

- Disable NAT Gateway in non-prod environments
- Use VPC endpoints for S3 and DynamoDB access
- Right-size compute for non-prod workloads

Estimates are directional. They do not affect billing. They inform decisions before billing happens.

What platform teams ship with DevX

H2: Cost gates engineers actually want to use.

PR cost diff Every infra PR gets an inline cost impact estimate before reviewers see it. The engineer sees what their change costs. Their reviewer sees it too. The conversation happens in the PR, not in a cost review meeting three weeks later.

Policy-as-code Cost guardrails versioned in your repo. Applied as advisory or required checks. Advisory mode: the PR gets an estimate, the engineer decides. Required mode: PRs over a defined cost threshold need explicit approval before merge. Both modes live in your repo as code: version controlled, reviewable, auditable.

Native CI integration GitHub Actions, GitLab CI, Azure Pipelines, Jenkins, Argo. No new pipeline required. DevX slots into what your teams already use.

Multi-IaC support Terraform, OpenTofu, Pulumi, CloudFormation, Helm, Kubernetes, raw Kubernetes manifests. 7+ formats supported.

Who DevX is for

Platform engineers Stop cost governance from being a retrospective exercise. Policies live in the repo. Every PR gets a signal before it ships. Governance happens where the work happens.

Application engineers Catch costly code patterns early. Expensive loops, chatty APIs, and inefficient resource usage flagged in context before production. The estimate arrives with a suggested fix.

FinOps teams Shift cost accountability into the delivery workflow. Surface cost risks where decisions are made, before infrastructure or code ships. Stop chasing post-production waste.

Data teams Detect expensive queries, inefficient scans, and over-provisioned compute in PRs and CI jobs. Catch the problem before the data platform gets blamed.

Pricing

H2: Most customers recover the cost of DevX from a single prevented regression.

A missed NAT gateway cleanup runs \$800 per month minimum. A misconfigured always-on instance in non-prod runs higher. DevX catches these before they merge.

Free — \$0/month

- Unlimited public repos
- Standard PR cost estimates
- Basic CI scan
- Weekly reports

Business — \$900/month

- Everything in Free

- Private repos
- Priority CI scans
- Advanced PR policy engine
- SSO (GitHub, Okta, and others)
- Includes 1,000 units per month. \$0.50 per additional unit.

Enterprise — Custom

- Everything in Business
 - Custom allowance and unlimited scale
 - Dedicated success manager
 - SLA and premium support
 - On-prem deployment options
-

FAQ

Will this slow down our PR flow? No. DevX runs as a fast check and posts a single inline comment. No extra approval step unless you opt into required mode.

Which IaC tools are supported? Terraform, OpenTofu, Pulumi, CloudFormation, Helm, Kubernetes, raw Kubernetes.

How do we manage cost policy? Policies are code in your repo. Define thresholds, apply advisory or required enforcement per team or repo. Version controlled and reviewable like any other config.

What is the path from PR comment to actual savings? The engineer sees the estimate, adjusts the change before merge, and the regression never reaches production. The saving is the cost of the avoided change multiplied by its lifetime.

Final CTA

H2: See DevX catch a cost regression in your repo.

Connect your first account in under 30 minutes. Most teams find their first cost regression the same day.

CTA: Book a Demo **CTA Secondary:** Talk to Sales

H1: Multi-Cloud Cost Intelligence for Every Team

Subhead: See every dollar by team, product, and provider. One model that reconciles to finance and explains itself to engineering.

CTA Primary: Book a Demo **CTA Secondary:** Explore the Platform

Stats strip

- \$738,983 recovered — Berkshire Hathaway HomeServices
 - \$101,736 saved annually
 - \$61,582 recovered in a single month
 - 5+ cloud providers unified
-

The problem

H2: Most FinOps tools show dashboards. This one changes what happens next.

Most organisations already have a FinOps tool. Most of those tools show dashboards that look correct and generate recommendations that go unactioned. The problem is not visibility. It is that the intelligence stops at the dashboard.

CloudVerse FinOps Platform connects spend data to the teams and decisions that drive it. What the bill says, what changed, who owns it, and what to do next.

Four product surfaces, one unified view

H2: Four product surfaces. One unified view.

Workload mapping Allocation that sticks. Every tag, account, BU, and shared-service split reconciled to the finance model. Teams see the spend they own. Leaders have a view that matches the finance report.

Anomaly detection Identify the team, environment, and charge that drove the anomaly within hours of it happening, not in the next billing cycle review. The alert arrives with attribution, not just a number.

Commitments with payback proof RI, SP, and CSP coverage modelled with explicit payback before you commit. Commitment planning that shows the math, not just the recommendation. No surprises.

Audit-ready chargeback Showback and chargeback wired to BU reporting with multi-currency and tax support. The finance team gets a report they can use. The engineering team gets a model they can trust.

Case study

H2: How Berkshire Hathaway HomeServices recovered \$738,983

\$101,736 recovered annually **\$61,582** recovered in a single month **\$738,984** total recovery

A growing mid-market cloud and AWS environment. Fragmented tagging, no team-level attribution, and a FinOps process that depended on manual variance reviews.

CloudVerse FinOps Platform connected spend to teams, surfaced the highest-impact anomalies, and gave the team an allocation model that held up to finance scrutiny. The waste was always there. It just had no address.

How it works

H2: Inform. Optimise. Operate.

Inform Allocation, tagging, attribution, and variance reporting. Understand what is happening, who owns it, and what changed since last month.

Optimise Rightsizing, commitment coverage, and waste identification. Prioritised by impact. Engineering-grade signals alongside finance reporting.

Operate Automated recommendations, chargeback, budget alerts, and governance controls. Automation paths with approval workflows and full audit logs.

Built for the teams accountable for the bill

FinOps Manager One model that covers cloud, data platforms, and AI workloads. Everything reconciles to finance. Automation reduces the manual review cycle.

Cloud Engineer Clear cost signals without leaving your workflow. See what changed, what it costs, and what to do. No more cost governance as a separate process.

CFO / VP Finance Spend accountability at team and product level. The bill explained in terms of business decisions, not infrastructure line items.

FAQ

How is this different from cost explorer dashboards? Dashboards explain the invoice. CloudVerse governs the decisions that shape it: PR-level checks, AI/GPU economics, and warehouse query attribution on one model.

Can we keep our current allocation logic? Yes. CloudVerse imports existing tag structures and cost allocation rules as a starting point.

How long until we are using it day to day? Most teams have a working allocation model within the first account connection. Meaningful depth takes two to four weeks of iteration depending on tagging maturity.

Do you support multi-currency and tax? Yes.

Connects to the stack

AWS, Azure, Google Cloud, plus AIX for AI workload economics, DataX for warehouse intelligence, and DevX for shift-left cost intelligence.

Final CTA

H2: Bring your FinOps program onto one control plane.

Connect your first account in under 30 minutes. Most teams have their first non-obvious finding the same day.

CTA: Book a Demo **CTA Secondary:** Talk to Sales

PAGE 6: Solutions — AI Engineering (solutions/ai-engineering)

Breadcrumb: For AI Engineering

Hero

H1: Run Every AI Workload Where It Costs Least and Runs Best

Subhead: Live cost-quality routing across 8+ GPU and LLM providers. Policy-bound at execution. FinOps attribution included.

CTA Primary: Book a Demo **CTA Secondary:** Explore the Platform

Stats strip

- 8+ GPU and LLM providers
 - Live cost-quality routing
 - Policy-bound at execution
 - Auditable decision trail
-

The situation AI engineering teams are in

Most teams hardcoded a model endpoint eighteen months ago because it was the right call at the time. Since then, three providers have launched models that handle 70% of those requests at a third of the cost. The endpoint still works, so nobody has changed it.

Meanwhile, GPU costs for agent infrastructure vary by 3x month to month depending on workload activity. Finance asks for an allocation report. The answer is a spreadsheet that took two engineers a week to build.

AIX solves both of these.

What AIX unlocks for AI engineering teams

Multi-provider routing Score every request across OpenAI, Anthropic, Bedrock, Vertex, Cohere, Groq, HuggingFace, and self-hosted infrastructure. Route to the best fit based on cost, latency, and quality requirements for that specific task type. Automatically.

Policy guardrails Compliance, content, latency, and budget guardrails enforced before execution. PII handling rules, data residency constraints, and provider allowlists applied at the routing layer, not retrofitted in application code.

Right-sized GPU economics Move workloads between hosted and dedicated GPU pools without rewriting code. AIX handles the routing logic. Your engineers handle the model.

Token and GPU attribution Spend tied back to team, product, and unit revenue. Automatic. No separate AI cost report reconstructed from billing data after the fact.

Model registry Track model versions, provider options, and performance benchmarks. Multi-region failover built in. Primary and fallback routes generated on every request.

The numbers

- 40-90% cost reduction across production workloads
 - Less than 15ms routing overhead
 - 96.8% lower cost vs hardcoded Claude Sonnet
 - Spend variance drops from 3x to under 15% for agent workloads after AIX
-

Platform that powers this solution

AIX: GPU and LLM economics **FinOps Platform:** Multi-cloud cost intelligence **DevX:** Shift-left cost intelligence

AI engineering questions answered

How is this different from an LLM gateway? A gateway executes a routing rule you wrote. AIX evaluates whether that rule is still the right one, continuously, using live cost and latency signals. The distinction matters when pricing shifts or a new model becomes available.

Can we mix hosted and self-hosted models? Yes. AIX treats private GPU deployments and NeoCloud infrastructure as first-class routing targets alongside managed APIs. Same cost and policy logic applies to all of them.

What about prompt privacy? AIX does not proxy traffic by default. It returns a routing decision. Your application sends the request. Prompt content never passes through AIX unless you configure otherwise.

How does the cost attribution work? Every request is tagged to the team, feature, or workload that made it. Cost allocation is automatic. Finance does not need engineering to explain the report.

PAGE 7: Solutions — Data Teams (solutions/data-teams)

Breadcrumb: For Data Teams

Hero

H1: Find the Queries Quietly Running Up Your Bill

Subhead: Query, dashboard, and dbt-model-level attribution across Snowflake, Databricks, BigQuery, Microsoft Fabric, and Synapse. Safe automation when you want it.

CTA Primary: Book a Demo **CTA Secondary:** Explore the Platform

Stats strip

- \$117.16 — full scan detected
 - 334.6 GB — scanned per query
 - 6 warehouses supported
 - Read-only by default
-

The situation data teams are in

Snowflake costs went up again. The data team got a Slack message asking what happened. Forty minutes of log diving later, someone found a query running 334 GB of full table scans 77 times a month. No partition filter. \$117 in one query pattern. Nobody knew.

DataX finds these automatically. Attribution down to the SQL. The fix included.

What data teams unlock with DataX

Query attribution Every query tied to the user, role, dashboard, model, or job that ran it. When the data team gets blamed for the bill, they can show exactly which workload and which team owns the cost.

Pattern detection Repeated expensive patterns surfaced with rewrite suggestions. DataX groups queries into cost-amplifying patterns, not one-off executions. The problem is usually a handful of patterns running hundreds of times.

Predictive signals Predict warehouse cost spikes from queue depth and pattern shifts before they hit the bill. Most regressions are visible in telemetry before they become a finance conversation.

Safe automation One-click partition, cluster, and right-size fixes. Policy-bound. Auditable. Reversible. Automation in DataX is opt-in, scoped, and logged in full.

The \$117 finding (real DataX output)

A BigQuery SELECT scanning 334.6 GB per query due to missing partition pruning.

- 0% cache hit rate
- 77 runs at \$1.52 average cost
- \$117.16 total
- Tagged: full-scan, spiky, fan-out
- Fix: one-click partition pruning

This is where hidden spend actually lives. Not in one catastrophic bill. In a pattern nobody noticed until the month-end review.

Platform that powers this solution

DataX: Warehouse intelligence **FinOps Platform:** Multi-cloud cost intelligence **DevX:** Shift-left cost intelligence

Data team questions answered

Does DataX read our actual data? No. DataX reads query history and metadata only, never row contents. Read-only by default.

Which warehouses are supported? Snowflake, Databricks, BigQuery, Microsoft Fabric, Azure Synapse.

How does it work with dbt? DataX attributes costs to dbt model runs and flags high-cost models with optimisation suggestions. Attribution follows the DAG.

Can data engineers act on recommendations directly? Yes. One-click optimisation. All actions are policy-bound and auditable. No finance approval cycle required for routine fixes.

PAGE 8: Solutions — FinOps Teams (solutions/finops-teams)

Breadcrumb: For FinOps Teams

Hero

H1: The Control Plane Your Finance and Engineering Teams Both Trust

Subhead: Allocation, anomaly response, commitments, and chargeback on one model. Reconciles to finance. Explains itself to engineering.

CTA Primary: Book a Demo **CTA Secondary:** Explore the Platform

Stats strip

- \$738,983 recovered (BHHS)
 - First account connected in less than 30 minutes
 - Same day first finding
 - 5+ cloud providers unified
-

The situation FinOps teams are in

The dashboards are accurate. The recommendations are ignored. Engineering says the tool does not understand their workloads. Finance says the numbers do not match the invoice. Leadership wants automation. You are running manual variance reports to bridge the gap.

CloudVerse gives you one model that covers cloud infrastructure, data platforms, and AI workloads. Attribution that holds up to finance scrutiny. Recommendations that engineering actually acts on.

What FinOps teams ship faster with CloudVerse

Allocation everyone agrees on One model with tags, accounts, BUs, and shared-service splits,

reconciled to finance. The finance team gets the report they need. The engineering team gets a view that matches their cost centres.

Anomalies with attribution Identify the team, environment, and charge that drove the anomaly within hours of it happening. The alert arrives with context, not just a number. No more spending half the week explaining what spiked.

Commitments with payback proof RI, SP, and CSP coverage modelled with explicit payback before you commit. The decision gets made with the full math visible, not retrospectively justified.

Audit-ready chargeback Showback and chargeback wired to BU reporting with multi-currency and tax support. The report your finance team can actually use.

Case study

H2: How Berkshire Hathaway HomeServices recovered \$738,983

\$101,736 annual recovery. \$61,582 in a single month. \$738,984 total.

Fragmented tagging, no team-level attribution, and a FinOps process that depended on manual variance reviews. CloudVerse connected spend to teams, surfaced the highest-impact anomalies, and delivered an allocation model the finance team trusted.

How this is different

Dashboards explain the invoice. CloudVerse governs the decisions that shape it: PR-level checks in engineering workflows, AI and GPU economics, and warehouse query attribution. All on one model.

Platform that powers this solution

FinOps Platform: Multi-cloud cost intelligence **DevX:** Shift-left cost intelligence **DataX:** Warehouse intelligence

FinOps team questions answered

How is this different from cost-explorer dashboards? Dashboards explain the invoice. CloudVerse governs the decisions that shape it.

Can we keep our current allocation logic? Yes. CloudVerse imports existing tag structures and cost allocation rules as a starting point.

How long until we are using it day to day? Most teams have a working allocation model within the first account connection.

Do you support multi-currency and tax? Yes.

PAGE 9: Solutions — Platform Engineering (solutions/platform-eng)

Breadcrumb: For Platform Engineering

Hero

H1: Cost Gates Engineers Actually Want to Use

Subhead: PR-level cost diffs, policy-as-code, and right-sizing recommendations in your IaC pipeline. Quantified savings rolled into FinOps.

CTA Primary: Book a Demo **CTA Secondary:** Explore the Platform

Stats strip

- +\$2.4k to \$2.9k flagged from a single PR
 - Free tier available
 - GitHub, GitLab, Azure DevOps native
 - 7+ IaC formats supported
-

The situation platform engineering teams are in

You have tried sending cost reports to engineering teams. The response is silence or confusion. The numbers do not connect to anything engineers can act on at the moment they are making decisions.

DevX puts the cost estimate in the PR. Before the code ships. While the engineer still has context on what they built and why.

The cost comment arrives the same way a test failure does. In context. With a suggested fix. Before the merge.

What platform teams ship with CloudVerse

PR cost diff Every infra PR gets an inline cost impact estimate before reviewers see it. Advisory or required, your choice. The governance happens where the work happens.

Policy-as-code Cost guardrails versioned in your repo. Advisory mode posts an estimate. Required mode blocks PRs above your defined cost threshold until explicitly approved. Both modes are code, not configuration buried in a third-party portal.

Native CI integration GitHub Actions, GitLab CI, Azure Pipelines, Jenkins, Argo. No new pipeline. DevX slots into what your teams already use.

Multi-IaC support Terraform, OpenTofu, Pulumi, CloudFormation, Helm, Kubernetes. Seven-plus formats. One integration.

Platform that powers this solution

DevX: Shift-left cost intelligence **FinOps Platform:** Multi-cloud cost intelligence **AIX:** GPU and LLM economics

Platform engineering questions answered

Will this slow down our PR flow? No. DevX runs as a fast check and posts a single inline comment. No extra approval step unless you opt into required mode.

Which IaC tools are supported? Terraform, OpenTofu, Pulumi, CloudFormation, Helm, Kubernetes, raw Kubernetes.

How do we manage cost policy? Policies are code in your repo. Version controlled, reviewable, auditable like any other config.

What is the path from PR comment to actual savings? The engineer adjusts the change, the regression never reaches production. The saving is the cost of the avoided change multiplied by its lifetime.

PAGE 10: Solutions — Enterprise (solutions/enterprise)

Breadcrumb: For Enterprise

Hero

H1: One Control Plane for Multi-Cloud, AI Infrastructure, and Data Economics

Subhead: The compute economics platform global enterprises run on. SSO, SCIM provisioning, granular RBAC, audit logs, customer-managed encryption keys, and the procurement options the office of the CIO expects.

CTA Primary: Talk to Our Enterprise Team **CTA Secondary:** Explore the Platform

Stats strip

- SSO + SCIM — identity-ready on day one

- US / EU / APAC — residency options
 - AWS, Azure, GCP — marketplace listed
 - 24/7 enterprise support
-

What enterprise teams operationalise on day one

One model across the estate Cloud, AI infrastructure, and warehouse spend on a single allocation model. One view for FinOps, one view for engineering, one view for finance. They all match.

Identity and audit SSO, SCIM, granular RBAC, audit logs, customer-managed encryption keys. The access controls and audit trails your security and compliance teams require, available from the start.

Regional residency US, EU, and APAC regions with private-link and VPC options. Data stays where your sovereignty requirements say it should.

Marketplace and procurement AWS, Azure, and Google Cloud Marketplace listings with committed-spend redemption. Procurement through the channels your finance team already uses.

Modules that power enterprise deployments

FinOps Platform: Multi-cloud cost intelligence **AIX:** GPU and LLM economics **DevX:** Shift-left cost intelligence

Enterprise questions answered

What enterprise controls are included? SSO/SAML, OIDC, SCIM provisioning, granular RBAC, audit logs, customer-managed encryption, and regional data residency. All included.

Where is data stored? Region-specific. US, EU, and APAC options available with private-link and VPC options.

Do you support air-gapped or VPC deployments? Contact our enterprise team to discuss your specific requirements.

What about procurement? AWS, Azure, and Google Cloud Marketplace with committed-spend redemption.

Who do we work with for rollout? CloudVerse enterprise accounts include a dedicated solutions engineer for onboarding and a customer success manager post-launch.

Final CTA

H2: Talk to our enterprise team.

Connect your first account in under 30 minutes. Most teams have their first non-obvious finding the same day.

CTA: Book a Demo **CTA Secondary:** Talk to Sales

PAGE 11: Integrations (integrations)

Hero

H1: Connects to the Stack Your Teams Already Use

Subhead: Add your integration once. CloudVerse handles routing, attribution, and cost tracking across all of them. No code changes when you add a new provider.

CTA: Set up your integrations

Security model (lead with this)

H2: Read-only by default. Automation is opt-in.

CloudVerse connects to your infrastructure using read-only access. It reads metadata, query logs, billing telemetry, and policy signals. It never touches your underlying data, workload code, or runtime configuration unless you explicitly grant automation permissions.

Automation permissions are scoped, explicit, and auditable. Every automated action is logged with the reason, the expected impact, and a rollback path.

Permissions model:

- Read-only by default: metadata, query logs, billing telemetry. Never touches underlying data.
 - Automation is opt-in: explicitly granted, scoped to specific resources, fully auditable.
 - Least-privilege scripts: provided for your security team to review before connection.
-

Integration categories

Cloud providers AWS, Microsoft Azure, Google Cloud Platform

Data platforms Snowflake, Databricks, BigQuery, Microsoft Fabric, Azure Synapse

AI providers OpenAI, Anthropic, HuggingFace, Mistral AI, Llama/Ollama, Cohere, Together AI, Groq, DeepSeek

Code and CI GitHub, GitLab, Azure DevOps, GitHub Actions, GitLab CI, Jenkins, Argo

APM and telemetry OpenTelemetry, Datadog, New Relic, Dynatrace

Identity Okta, Azure Active Directory, Google Workspace (SSO + SCIM)

AI provider detail

H2: Supported model providers

Provider	Models and notes
OpenAI	GPT-4o, GPT-4o-mini, GPT-4 Turbo, GPT-3.5 Turbo
Anthropic	Claude 3.5 Sonnet, Claude 3 Haiku, Claude 3 Opus
HuggingFace	Hosted inference endpoints and open models via HuggingFace Hub
Mistral AI	Mistral Large, Mistral Small, Mixtral 8x7B
Llama / Ollama	Meta Llama 3 family via Ollama and compatible deployments
Cohere	Command R, Command R+, reranking endpoints
Together AI	Open model inference at scale
Groq	Ultra-low latency LPU inference
DeepSeek	DeepSeek-V3, DeepSeek-R1

More providers added regularly.

GPU infrastructure

H2: Private deployments and GPU infrastructure.

AIX treats private GPU capacity as a first-class routing target alongside managed APIs. If you run models on dedicated hardware or a NeoCloud provider, AIX routes to them with the same cost and policy logic.

- Private model deployments: vLLM, TGI, custom inference
 - GPU providers: CoreWeave, Lambda Labs, RunPod
 - On-premises infrastructure: your own inference hardware
-

Enterprise access patterns

H2: Built for enterprise access patterns.

Private-link available: No data leaves your VPC. **Audit logs:** All connection and access activity logged. **Customer-managed encryption keys:** Available for enterprise accounts. **Least-privilege onboarding scripts:** Reviewed by your security team before any connection is established. **Multiple ingestion modes:** Push, pull, and event-driven options depending on platform type.

Three steps to get started

1. Connect providers: Add API keys and endpoints. Minutes per provider.
 2. Define constraints: Set latency, budget, compliance, and provider rules per workload.
 3. Route through CloudVerse: Point model calls at the CloudVerse endpoint. Routing handled from there.
-

Final CTA

H2: Need a connector that is not listed?

Talk to us. We add integrations based on customer demand.

CTA: Talk to Us **CTA Secondary:** View Documentation

Content confidence score: 91/100 KB relevance: 95% / Competitor positioning improvement: 93% / Above-industry-average probability: 88% / Output quality: 89% Weighted average (0.3 / 0.25 / 0.25 / 0.2): 91